

Audio-identification providers — evaluation for SoundtrackID

Evaluated 2026-08-25. Use case: identify short film-music cues and return title / artist / album and — critically — the ISRC, which we chain via MusicBrainz → work → composer/lyricist → death year for the public-domain assessment.

The decisive question

Not "does it fingerprint audio?" but **"does it ship with a catalog of commercial music to match against?"** Two very different product categories:

- **Catalog-backed recognition services** — they hold the music database, so they can identify an *unknown* cue. The only ones usable for us.
- **Self-hosted ACR engines** — you must build/maintain your *own* reference fingerprint database. Great for ad/broadcast monitoring or detecting repeats of your own content; useless for identifying arbitrary commercial film music we don't already own.

Comparison

Service	Type	Music catalog?	Short-cue recognition	ISRC / metadata	Access & cost	Verdict
ACRCloud	Commercial API	Yes, huge	Yes (snippet fingerprint)	ISRC + Spotify/YT/MusicBrainz	REST, paid (trial tier)	✅ Best — current default
AudD	Commercial API	Yes, 160M+	Yes (12s chunks; enterprise = full file + timestamps)	ISRC, UPC, label (Startup+)	REST, 300 free then \$5→\$2 / 1k	✅ Strong alternative / cross-check
Shazam (ShazamKit)	Apple framework	Yes, full Shazam catalog	Excellent	ISRC + Apple Music id	Free w/ Apple Dev; native Swift framework, not REST	✅ Usable (we use unofficial shazamio)
Gracenote	Commercial API	Yes, huge	Yes	ISRC (must request in query)	Dev license non-commercial only; commercial = enterprise sales	⚠️ Maybe — heavy onboarding/licensing
AcoustID	Open / MusicBrainz	Yes (via MusicBrainz)	No — needs ~full track (exact-file fingerprint)	MBID → ISRC/composer via MB	Free / open	⚠️ Limited — weak on short cues
EmySound	Self-host ACR engine	No (bring-your-own DB)	n/a	No	Free (non-commercial) / self-host	❌ No commercial catalog
Veritone (aiWARE)	Enterprise ACR platform	No (your libraries)	n/a	No	Enterprise	❌ Wrong tool
Audioneex	Open-source engine (MPL-2.0)	No (bring-your-own DB)	n/a	No	Free/open + commercial SDK	❌ Engine only; no catalog
MADCAT (Oxford Wave)	Forensic dedup/sync	No	n/a — matches overlapping content between your recordings	No	Commercial	❌ Not a music-ID service

Recommendation

- **Keep ACRCloud as primary** — best fit, already integrated, returns ISRC.
- **Add AudD as a secondary / cross-check** — cleanest drop-in: simple REST, cheap, returns ISRC + MusicBrainz links. Good for A/B-ing against ACRCloud to cut "later-sample" false matches.
- **Shazam** stays as secondary; the only upgrade is going native ShazamKit for reliability/legitimacy vs. the unofficial wrapper.

- **AcoustID** only as a free fallback for longer segments (its full-file fingerprint needs the length), since it chains straight into our MusicBrainz→composer→PD logic.
- **EmySound / Audioneex / Veritone / MADCAT** — not usable: they match against a reference DB you supply. (Only relevant if we later want to detect repeats of the same cue within one film — self-fingerprinting — a different feature.)

Sources

audd.io · emysound.com · audioneex.com/licensing · oxfordwaveresearch.com/products/madcat · veritone.com/aiware/ai-engines/audio-fingerprinting · developer.gracenote.com/isrc-code · developer.apple.com/documentation/ShazamKit · en.wikipedia.org/wiki/AcoustID